



CSE 180, Database Systems I

Shel Finkelstein, UC Santa Cruz

Lecture 5

Chapter 4 : Intermediate SQL, Part 1

Database System Concepts, 7th Edition.
©Silberschatz, Korth and Sudarshan
Modified by Shel Finkelstein for CSE 180



Important Notices

- **Reminder:** Midterm is on **Wednesday, February 12.**
 - No make-ups, no early Midterms, no late Midterms ... and no devices.
 - You may bring a **single two-sided 8.5" x 11" sheet of paper** with as much info written (or printed) on it as you can fit and read unassisted.
 - No sharing of these sheets will be permitted.
 - “Practice Midterm” from Fall 2019 has been posted on Piazza under Resources→Exams.
 - We’ll be covering more material before Midterm this quarter.
 - Solution has also been posted under Resources→Exams
 - ... but I strongly suggest that you take it yourself first, rather than just reading the solution.
 - Hope that all requests for DRC accommodation have been submitted.
 - DRC students will receive access to webcasts only if that specific accommodation has been requested by DRC.
 - Piazza announcement about Midterm describes required seating pattern for Midterm.



Important Notices

- Lab2 solution has been posted on Piazza as lab2.zip under Resources→Lab2.
- Lab3 has been posted on Piazza, with a lab3_create.sql file that creates tables for Lab3.
 - A lab3_data_loading.sql file that is **needed** to do Lab3 will be posted soon.
 - This Lecture, Lecture 5, includes material for Lab3.
 - A transaction (from Lecture 6) is also used in Lab3.
 - Lab3 is due on **Sunday, Feb 23, 11:59pm**. You have 3 weeks to complete Lab3 because the CSE 180 Midterm is on Wednesday, Feb 12.
- The third Gradiance Assignment, "CSE 180 Winter 2020 #3", will be assigned after the Midterm.
- Access to all webcasts went away on Wednesday, January 29.
 - Attend Lectures, Lab Sections, Office Hours and LSS Tutoring.
- See [Piazza notices](#) about LSS Tutoring with Jeshwanth Bheemanpally.



Chapter 4: Intermediate SQL

- Views
- Indexes
- Integrity Constraints
- Transactions
- Join Expressions
- SQL Data Types and Schemas
- Authorization



Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructor's name and department, but not the salary. This person should see a relation described, in SQL, by

```
SELECT ID, name, dept_name  
FROM instructor;
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.



Some Advantages of Views

- **Short-hand/encapsulation:** You can treat a view as a table (for queries), which allows you to define a short-hand for a concept that might involve a complicated query.
- **Re-use:** You can re-use a view as often as you like; other people who can access the view may also re-use it.
- **Authorization:** People may be granted access to a view, even though they don't have access to the underlying tables.
 - They may not even know that the view isn't a table.
 - Even if they know that it's a view, they may not know which tables underlie it.
- **Logical data independence:** Even if the tables underlying a view change, the view may still be used in queries (or applications), without re-writing the queries (or applications).



View Definition

- A view is defined using the **CREATE VIEW** statement which has the form

CREATE VIEW *v* AS <query expression>

where <query expression> is any legal SQL query.

- The view name is *v*.
- After a view has been defined, the view name can be used to refer to the virtual relation corresponding to the view.
- View definition does not create a new relation!
 - Rather, creating a view saves the expression defining the view in the DBMS's system catalog.
 - The DBMS automatically substitutes that expression into any SQL statements that use the view.



View Definition and Use

- Define a view of instructors that doesn't give their salaries.

```
CREATE VIEW faculty AS  
    SELECT ID, name, dept_name  
        FROM instructor;
```

- Find all instructors in the Biology department using the view.

faculty

```
SELECT name  
FROM faculty  
WHERE dept_name = 'Biology';
```

- Create a view that provides total salary for each department.

```
CREATE VIEW departmentTotalSalary(dept_name, total_salary) AS  
    SELECT dept_name, SUM(salary)  
        FROM instructor  
        GROUP BY dept_name;
```

or

```
CREATE VIEW departmentTotalSalary AS  
    SELECT dept_name, SUM(salary) AS total_salary  
        FROM instructor  
        GROUP BY dept_name;
```



Views Defined Using Other Views

Views may be used (together with other views and tables) to define another view.

- **CREATE VIEW** *physics_fall_2017* **AS**
SELECT *C.course_id, S.sec_id, S.building, S.room_number*
FROM *course C, section S*
WHERE *C.course_id = S.course_id*
 and *C.dept_name = 'Physics'*
 and *S.semester = 'Fall'*
 and *S.year = '2017';*
- **CREATE VIEW** *physics_fall_2017_watson* **AS**
SELECT *course_id, room_number*
FROM *physics_fall_2017*
WHERE *building = 'Watson';*



Evaluating a View Defined on a View

- The view:

```
CREATE VIEW physics_fall_2017_watson AS
    SELECT course_id, room_number
    FROM physics_fall_2017
    WHERE building = 'Watson';
```

- Becomes:

```
CREATE VIEW physics_fall_2017_watson AS
    SELECT P.course_id, P.room_number
    FROM ( SELECT C.course_id, S.building, S.room_number
            FROM course C, section S
            WHERE C.course_id = S.course_id
                  AND C.dept_name = 'Physics'
                  AND S.semester = 'Fall'
                  AND S.year = '2017' ) AS P
    WHERE P.building = 'Watson';
```



Evaluating a Query Defined on a View

- The query:

```
SELECT P.course_id, P.room_number, Teach.ID  
FROM physics_fall_2017 P, teaches Teach  
WHERE P.course_id = Teach.course_ID  
    AND Teach.semester = 'Fall'  
    AND Teach.year = '2017';
```

- Becomes:

```
SELECT P.course_id, P.room_number, Teach.ID  
FROM ( SELECT C.course_id, S.building, S.room_number  
        FROM course C, section S  
        WHERE C.course_id = S.course_id  
            AND C.dept_name = 'Physics'  
            AND S.semester = 'Fall'  
            AND S.year = '2017' ) AS P,  
    teaches Teach  
WHERE P.course_id = Teach.course_ID  
    AND Teach.semester = 'Fall'  
    AND Teach.year = '2017';
```



DROP VIEW

```
CREATE VIEW physics_fall_2017 AS
  SELECT C.course_id, S.sec_id, S.building, S.room_number
  FROM course C, section S
  WHERE C.course_id = S.course_id
        and C.dept_name = 'Physics'
        and S.semester = 'Fall'
        and S.year = '2017';
```

```
DROP VIEW physics_fall_2017 ;
```

- What happens if after DROP VIEW you execute the following?
 - **SELECT * FROM** physics_fall_2017;
 - **SELECT * FROM** course;
 - **SELECT * FROM** section;
 - **SELECT * FROM** physics_fall_2017_watson;



Materialized Views

- Normally, only the view definition is stored in the database; the contents of the view are not stored.
- Certain database systems allow some view relations to be physically stored, if explicitly requested.
 - A physical copy of the contents is stored when the view is defined.
 - Such views are called **Materialized Views**.
- If relations used in the query defining the view are updated, then the Materialized View contents becomes out-of-date.
 - **Maintaining** the Materialized View would require changing the view contents whenever its underlying relations are updated.
 - This can be expensive, so it's unusual to maintain Materialized Views.
 - Moreover, some views are not updatable ... as we'll see.



“Update” of a View

- Add a new tuple to *faculty* view which we defined earlier

INSERT INTO *faculty*

VALUES ('30765', 'Green', 'Music');

- This insertion must be represented by the insertion into the *instructor* relation
 - ... so it must have a value for salary.
- Two approaches
 - Reject the insert.
 - Insert the following tuple into the *instructor* relation:
('30765', 'Green', 'Music', **NULL**)
- Which approach is used?
 - What if salary can't be **NULL**?
 - What if salary can be **NULL**?



But Some “Updates” Can’t be Translated!

- **CREATE VIEW** *instructor_info* **as**
SELECT *instructor.ID*, *instructor.name*, *department.building*
FROM *instructor*, *department*
WHERE *instructor.dept_name*= *department.dept_name*;
- **INSERT INTO** *instructor_info*
VALUES ('69987', 'White', 'Taylor');
- What tuples should be inserted to make this work?
 - Easy ... insert tuple into *instructor* that has the *dept_name* of the department whose building is Taylor!
 - But which department, if multiple departments are in Taylor?
 - And what if no department is in Taylor?



View Updates in SQL

- Most SQL implementations allow updates only on very simple views.
 - The **FROM** clause has only one database relation.
 - The **SELECT** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **DISTINCT** specification.
 - Any attribute not listed in the **SELECT** clause can be set to **NULL**.
 - The query does not have a **GROUP BY** or **HAVING** clause.



Index Creation

- Many queries reference only a small proportion of the records in a table.
- It is inefficient for the system to read every record to find a record with particular value
- An **index** on an attribute of a relation is a data structure that allows the database system to find those tuples in the relation that have a specified value for that attribute efficiently, without scanning through all the tuples of the relation.
- We create an index with the **CREATE INDEX** command

CREATE INDEX <name> ON <relation-name> (attributes);



Index Creation Example

- **CREATE TABLE** *student*
(**ID** **VARCHAR** (5),
name **VARCHAR** (20) **NOT NULL**,
dept_name **VARCHAR** (20),
tot_cred **NUMERIC** (3,0) **DEFAULT** 0,
PRIMARY KEY (**ID**))
- **CREATE INDEX** *studentID_index* **ON** *student*(**ID**)
- The DBMS can execute the query:

```
SELECT *  
FROM student  
WHERE ID = '12345'
```

by using the index to find the required record, without looking at all records of *student*.



Table Modifications and Indexes

- If a table is updated, all Indexes on that table are immediately automatically updated within the same transaction.
 - Which indexes do you need to change on INSERT and DELETE?
 - What about UPDATE?
- SQL statements can be executed regardless of which indexes (if any) exist in the database.
 - Physical Independence!
 - Applications don't have to be modified when indexes are created or dropped!
- What gets impacted when indexes are created or dropped?
 - Performance of SQL statements
 - Some may run faster, some may run slower.



Benefits of Indexes for Queries

```
SELECT *  
FROM student  
WHERE name = 'Lee' AND tot_cred = 90;
```

How much would each of these indexes help?

```
CREATE INDEX Cred_Index ON student(tot_cred);
```

```
CREATE INDEX NameIndex ON student(name);
```

```
CREATE INDEX CNIndex ON student(tot_cred, name);
```

```
CREATE INDEX NCIndex ON student(name, tot_cred);
```

How much would each of the indexes help if the WHERE clause was just tot_cred = 90?



Benefits of Indexes for Another Query

```
SELECT *
FROM student
WHERE name = 'Lee' AND tot_cred < 90;
```

How much would each of these indexes help?

```
CREATE INDEX CredIndex ON student(tot_cred);
```

```
CREATE INDEX NameIndex ON student(name);
```

```
CREATE INDEX CNIndex ON student(tot_cred, name);
```

```
CREATE INDEX NCIndex ON student(name, tot_cred);
```

How much would each of the indexes help if the WHERE clause was just tot_cred < 90?



Disadvantages of Indexes?

- Why not put indexes on every attributes, or even on every combination of attributes that you might query on?
 - Huge number of indexes
 - Space for indexes
 - Cache impact of searching indexes
 - Update time for indexes when table is modified



Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - Every checking account must have a balance greater than \$100.00
 - The salary of a bank employee must be at least \$15.00 an hour.
 - A customer must have a (non-null) phone number.
 - Only known students can take classes.
 - If a student is taking a 5 credit course, then their GPA must be at least 3.0



Constraints on a Single Relation

- **NOT NULL**
- **PRIMARY KEY**
- **UNIQUE**
- **CHECK (P)**, where P is a predicate



NOT NULL Constraints

- **NOT NULL**

- Declare *name* and *budget* to be **NOT NULL**

name **VARCHAR(20) NOT NULL**

budget **NUMERIC(12,2) NOT NULL**



UNIQUE Constraints

- **UNIQUE** (A_1, A_2, \dots, A_m)
 - The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
 - Attributes in candidate keys may be **NULL** (unlike **PRIMARY KEY**).
 - If all the attributes A_1, A_2, \dots, A_m are not **NULL**, then there can't be multiple tuples that have the same value.
 - There may be multiple UNIQUE specifications for a table (unlike **PRIMARY KEY**).



CHECK Constraint

- A **CHECK** (P) constraint specifies that a predicate P that must be satisfied by every tuple in a relation.
- Example: Ensure that semester is one of fall, winter, spring or summer.

CREATE TABLE *section*

```
( course_id VARCHAR (8),  
sec_id VARCHAR (8),  
semester VARCHAR (6)  
      CHECK ( semester IN ('Fall', 'Winter', 'Spring', 'Summer') ),  
year NUMERIC (4,0),  
building VARCHAR (15),  
room_number VARCHAR (7),  
time slot id VARCHAR (4),  
PRIMARY KEY (course_id, sec_id, semester, year) );
```

- A **CHECK** constraint is violated if its value is FALSE. It is not violated if its value is TRUE or if its value is UNKNOWN.



Multiple Attributes in a CHECK Constraint

CHECK can refer to multiple attributes in the tuple. But if it does, then the **CHECK** must be a separate schema element.

- Ensure that if year is after 2019, then the building is not Watson.

CREATE TABLE section

```
( course_id VARCHAR (8),  
  sec_id VARCHAR (8),  
  semester VARCHAR (6),  
  year NUMERIC (4,0),  
  building VARCHAR (15),  
  room_number VARCHAR (7),  
  time_slot_id VARCHAR (4),  
  PRIMARY KEY (course_id, sec_id, semester, year),  
  CHECK ( year <= 2019 OR building <> 'Watson' ) );
```



Complex CHECK Constraints

- The SQL standard allows a predicate in a **CHECK** constraint to reference other tuples or relations, using a subquery.

Example:

```
CHECK ( time_slot_id IN  
        (SELECT time_slot_id FROM time_slot) )
```

- However, this capability is not implemented in most relational DBMS, so we won't discuss it.



Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If ‘Biology’ is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for ‘Biology’.
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the PRIMARY KEY of S. A is said to be a **FOREIGN KEY** of R if for any values of A appearing in R these values also appear in S.



Specifying Referential Integrity

- **Foreign keys** can be specified as part of the SQL **CREATE TABLE** statement
FOREIGN KEY (*dept_name*) **REFERENCES** *department*
 - By default, a foreign key references the primary-key attributes of the referenced table.
 - SQL allows a list of attributes of the referenced relation to be specified explicitly, so the attribute names don't have to be the same.
 - For example, if the key of *department* was *dname* instead of *dept_name*, we could write
- FOREIGN KEY** (*dept_name*)
REFERENCES *department* (*dname*)



Specifying Referential Integrity (Cont.)

If department's Primary Key was dname instead of dept_name, then we could write either of the following:

```
CREATE TABLE instructor (
    ID          char(5),
    name        VARCHAR(20) NOT NULL,
    dept_name   VARCHAR(20),
    salary      NUMERIC(8,2),
    PRIMARY KEY (ID),
    FOREIGN KEY (dept_name)
        REFERENCES department(dname);
```

```
CREATE TABLE instructor (
    ID          char(5) PRIMARY KEY,
    name        VARCHAR(20) NOT NULL,
    dept_name   VARCHAR(20) REFERENCES department(dname),
    salary      NUMERIC(8,2));
```



Specifying Referential Integrity (Cont.)

How would the **FOREIGN KEY** specification in *takes* change, if attribute names were different in *section*?

```
CREATE TABLE takes (
    ID          VARCHAR(5),
    course_id   VARCHAR(8),
    sec_id      VARCHAR(8),
    semester    VARCHAR(6),
    year        NUMERIC(4,0),
    grade       VARCHAR(2),
    PRIMARY KEY (ID, course_id, sec_id, semester, year) ,
    FOREIGN KEY (ID) REFERENCES student,
    FOREIGN KEY (course_id, sec_id, semester, year)
        REFERENCES section );
```

FOREIGN KEY (course_id, sec_id, semester, year)
REFERENCES section(course, sec, sem, year)



Actions to Ensure Referential Integrity

- When a referential-integrity constraint is violated, the default action is to reject the action that caused the violation.
- An alternative, in case of delete or update is to cascade

```
CREATE TABLE course (
  ...
  dept_name VARCHAR(20)
  FOREIGN KEY (dept_name) REFERENCES department
    ON DELETE CASCADE
    ON UPDATE CASCADE,
  . . .);
```

- Instead of **CASCADE** we can specify:
 - SET NULL**
 - SET DEFAULT**
- The specifications for **DELETE** and **UPDATE** may be different.
 - If either (or both) are omitted, default action (reject) is followed.

Argh; a Foreign Key can be NULL, so there won't be a corresponding value for it in the Referenced table!



Use of ALTER for Referential Integrity

- If instructor had been created without specifying dept_name as a Foreign Key (unusual but possible):

```
CREATE TABLE instructor (
    ID          char(5),
    name        VARCHAR(20) NOT NULL,
    dept_name   VARCHAR(20),
    salary      NUMERIC(8,2),
    PRIMARY KEY (ID) );
```

then we could do the following:

```
ALTER TABLE instructor
    ADD FOREIGN KEY (dept_name)
    REFERENCES department;
```

- What do you think would happen if an *instructor* tuple had a *dept_name* that isn't a *dept_name* in *department*?



Circular Integrity Constraints are a Problem!

- Consider:

```
CREATE TABLE person (
    ID char(10),
    name char(40),
    mother char(10) REFERENCES person,
    father char(10) REFERENCES person,
    PRIMARY KEY ID );
```

- How can we insert any tuples without causing a constraint violation?
 - Insert father and mother of a person before inserting person,
 - ... or set father and mother to **NULL** initially, update after inserting all persons (not possible if father and mother attributes are declared to be **NOT NULL**),
 - or **DEFER** constraint checking.



Assertions

- An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- The following constraints, can be expressed using assertions:
- For each tuple in the *student* relation, the value of the attribute *tot_cred* must equal the sum of credits of courses that the student has completed successfully.
- An instructor cannot teach in two different classrooms in a semester in the same time slot
- An assertion in SQL HAS the form:
CREATE ASSERTION <assertion-name> CHECK (<predicate>);
- Database systems do not implement assertions because they're too complicated and expensive.
 - We won't discuss them in this course.